

```
import pandas as pd
import numpy as np
```

```
df=pd.read_excel(r"D:\download\data (1).xlsx")
```

```
-----
-----
NameError                                Traceback (most recent call
last)
Cell In[1], line 1
----> 1 df=pd.read_excel(r"D:\download\data (1).xlsx")
```

```
NameError: name 'pd' is not defined
```

```
df.head(5)
```

	Income	Age	Dependents	Occupation	City_Tier
expenses \					
0	44637.249636	49	0	Self_Employed	Tier_1
1	26858.596592	34	2	Retired	Tier_2
2	50367.605084	35	1	Student	Tier_3
3	101455.600247	21	0	Self_Employed	Tier_3
4	24875.283548	52	4	Professional	Tier_2

	Loan_Repayment	Insurance	Groceries	Transport	...	\
0	0.000000	2206.490129	6658.768341	2636.970696	...	
1	0.000000	869.522617	2818.444460	1543.018778	...	
2	4612.103386	2201.800050	6313.222081	3221.396403	...	
3	6809.441427	4889.418087	14690.149363	7106.130005	...	
4	3112.609398	635.907170	3034.329665	1276.155163	...	

	Desired_Savings	Disposable_Income	Potential_Savings_Groceries	\
0	6200.537192	11265.627707	1685.696222	
1	1923.176434	9676.818733	540.306561	
2	7050.360422	13891.450624	1466.073984	
3	16694.965136	31617.953615	1875.932770	
4	1874.099434	6265.700532	788.953124	

	Potential_Savings_Transport	Potential_Savings_Eating_Out	\
0	328.895281	465.769172	
1	119.347139	141.866089	
2	473.549752	410.857129	
3	762.020789	1241.017448	
4	68.160766	61.712505	

Potential_Savings_Entertainment	Potential_Savings_Uilities	\
---------------------------------	----------------------------	---

0	195.151320	678.292859
1	234.131168	286.668408
2	459.965256	488.383423
3	320.190594	1389.815033
4	187.173750	194.117130

	Potential_Savings_Healthcare	Potential_Savings_Education \
0	67.682471	0.000000
1	6.603212	56.306874
2	7.290892	106.653597
3	193.502754	0.000000
4	47.294591	67.388120

	Potential_Savings_Miscellaneous
0	85.735517
1	97.388606
2	138.542422
3	296.041183
4	96.557076

[5 rows x 28 columns]

df.shape

(20000, 28)

df.columns

```
Index(['Income', 'Age', 'Dependents', 'Occupation', 'City_Tier',
      'expenses',
      'Loan_Repayment', 'Insurance', 'Groceries', 'Transport',
      'Eating_Out',
      'Entertainment', 'Utilities', 'Healthcare', 'Education',
      'Miscellaneous', 'Unnamed: 16', 'Desired_Savings_Percentage',
      'Desired_Savings', 'Disposable_Income',
      'Potential_Savings_Groceries',
      'Potential_Savings_Transport', 'Potential_Savings_Eating_Out',
      'Potential_Savings_Entertainment',
      'Potential_Savings_Utilities',
      'Potential_Savings_Healthcare', 'Potential_Savings_Education',
      'Potential_Savings_Miscellaneous'],
      dtype='object')
```

x=df.iloc[:,0:6].values

y=df.iloc[:,18].values

x

```
array([[44637.2496356859, 49, 0, 'Self_Employed', 'Tier_1',
       33371.621928834255],
```

```

[26858.5965917295, 34, 2, 'Retired', 'Tier_2',
17181.777859045094],
[50367.6050835768, 35, 1, 'Student', 'Tier_3',
36476.15445928645],
...,
[40604.5673726763, 30, 1, 'Professional', 'Tier_2',
38336.66223874504],
[118157.817239995, 27, 2, 'Professional', 'Tier_1',
107554.13242645186],
[8209.24976874268, 62, 3, 'Professional', 'Tier_1',
7348.899209965442]], dtype=object)

```

y

```

array([ 6200.53719244, 1923.1764339 , 7050.36042169, ...,
2267.90513393, 10603.68481354, 531.04400553])

```

```
df.isnull().sum()
```

```

Income      0
Age         0
Dependents  0
Occupation  0
City_Tier   0
expenses    0
Loan_Repayment  0
Insurance   0
Groceries   0
Transport   0
Eating_Out  0
Entertainment  0
Utilities   0
Healthcare  0
Education   0
Miscellaneous  0
Unnamed: 16    20000
Desired_Savings_Percentage  0
Desired_Savings  0
Disposable_Income  0
Potential_Savings_Groceries  0
Potential_Savings_Transport  0
Potential_Savings_Eating_Out  0
Potential_Savings_Entertainment  0
Potential_Savings_Uutilities  0
Potential_Savings_Healthcare  0
Potential_Savings_Education  0
Potential_Savings_Miscellaneous  0
dtype: int64

```

```
from sklearn.preprocessing import OrdinalEncoder
```

```

from sklearn.preprocessing import OneHotEncoder
from sklearn.compose import ColumnTransformer

from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=0)

transformer = ColumnTransformer(transformers=[
    ('encoder1', OrdinalEncoder(categories=[['Tier_3', 'Tier_2', 'Tier_1']]), [4]),
    ('encoder2', OneHotEncoder(drop="first"), [3])
], remainder='passthrough')

# Transform training and test data
X_train_transformed = transformer.fit_transform(x_train)
X_test_transformed = transformer.transform(x_test)

X_train_transformed
array([[2.0, 0.0, 1.0, ..., 36, 2, 34140.31020457392],
       [2.0, 0.0, 0.0, ..., 47, 4, 10244.395427778194],
       [1.0, 0.0, 0.0, ..., 35, 3, 98551.5712948812],
       ...,
       [0.0, 0.0, 0.0, ..., 31, 2, 15960.713927009849],
       [1.0, 0.0, 1.0, ..., 26, 4, 61929.80381889947],
       [2.0, 0.0, 0.0, ..., 43, 0, 37620.856361997394]], dtype=object)

from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train_transformed)
X_test_scaled = scaler.transform(X_test_transformed)

from sklearn.linear_model import LinearRegression
model = LinearRegression()
model.fit(X_train_scaled, y_train)

LinearRegression()

y_pred = model.predict(X_test_scaled)

y_pred
array([ -241.64897336,  5807.56395151, -138.19172317, ...,
        5907.43522445,  6493.56218864, 10084.07977449])

from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
mse = mean_squared_error(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
print("MSE:", mse)

```

```
print("MAE:", mae)
print("R2 Score:", r2)
```

```
MSE: 4648174.7459285725
MAE: 1407.369295311036
R2 Score: 0.9142419187040344
```

```
from sklearn.linear_model import RidgeCV, LassoCV
lasso = LassoCV(alphas=[0.01, 0.1, 1.0, 10.0], cv=5)
lasso.fit(X_train_scaled, y_train)
lasso_pred = lasso.predict(X_test_scaled)
```

```
print("\n Lasso Regression Results:")
print("Best alpha:", lasso.alpha_)
print("R2 Score:", r2_score(y_test, lasso_pred))
```

```
 Lasso Regression Results:
Best alpha: 0.01
R2 Score: 0.9142416336019124
```

```
import xgboost as xgb
```

```
model = xgb.XGBRegressor(n_estimators=100, learning_rate=0.1,
max_depth=3, reg_alpha=0, reg_lambda=1)
model.fit(X_train_scaled, y_train)
```

```
XGBRegressor(base_score=None, booster=None, callbacks=None,
             colsample_bylevel=None, colsample_bynode=None,
             colsample_bytree=None, device=None,
             early_stopping_rounds=None,
             enable_categorical=False, eval_metric=None,
             feature_types=None,
             feature_weights=None, gamma=None, grow_policy=None,
             importance_type=None, interaction_constraints=None,
             learning_rate=0.1, max_bin=None, max_cat_threshold=None,
             max_cat_to_onehot=None, max_delta_step=None, max_depth=3,
             max_leaves=None, min_child_weight=None, missing=nan,
             monotone_constraints=None, multi_strategy=None,
             n_estimators=100,
             n_jobs=None, num_parallel_tree=None, ...)
```

```
y_pred = model.predict(X_test_scaled)
r2 = r2_score(y_test, y_pred)
```

```
y_pred
r2
```

```
0.9148043147146206
```

```
y_pred
```

```
array([1069.4144, 5950.5566, 907.5372, ..., 5377.032 , 5762.297 ,
       7727.817 ], dtype=float32)
```

```
r2
```

```
0.9148043147146206
```

```
from sklearn.linear_model import Ridge, Lasso
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import mean_squared_error, r2_score
import numpy as np

# Define parameter grid
param_grid = {'alpha': [0.01, 0.1, 1, 10, 100]}

# ----- Ridge Regression -----
ridge = Ridge()
ridge_cv = GridSearchCV(ridge, param_grid, cv=5, scoring='r2')
ridge_cv.fit(X_train_scaled, y_train)

# Predict and evaluate Ridge
ridge_pred = ridge_cv.predict(X_test_scaled)
ridge_r2 = r2_score(y_test, ridge_pred)
ridge_rmse = np.sqrt(mean_squared_error(y_test, ridge_pred))

# ----- Lasso Regression -----
lasso = Lasso(max_iter=10000)
lasso_cv = GridSearchCV(lasso, param_grid, cv=5, scoring='r2')
lasso_cv.fit(X_train_scaled, y_train)

# Predict and evaluate Lasso
lasso_pred = lasso_cv.predict(X_test_scaled)
lasso_r2 = r2_score(y_test, lasso_pred)
lasso_rmse = np.sqrt(mean_squared_error(y_test, lasso_pred))

# ----- Print Results -----
print("\n Ridge Regression:")
print("Best alpha:", ridge_cv.best_params_['alpha'])
print("R2 Score:", ridge_r2)
print("RMSE:", ridge_rmse)

print("\n Lasso Regression:")
print("Best alpha:", lasso_cv.best_params_['alpha'])
print("R2 Score:", lasso_r2)
print("RMSE:", lasso_rmse)
```

```
\n Ridge Regression:
Best alpha: 1
R2 Score: 0.9142220221090633
RMSE: 2156.212689300464
```

□ Lasso Regression:

Best alpha: 0.1

R² Score: 0.9142390374803955

RMSE: 2155.998819842906

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
from statsmodels.stats.outliers_influence import variance_inflation_factor
from statsmodels.stats.stattools import durbin_watson
from statsmodels.stats.diagnostic import het_breuschpagan
from scipy.stats import shapiro
import statsmodels.api as sm

# Train model
model = LinearRegression()
model.fit(X_train_scaled, y_train)

# Predict
y_pred = model.predict(X_test_scaled)
residuals = y_test - y_pred

# 1. Linearity & Homoscedasticity (Residuals vs Predicted)
plt.figure(figsize=(6, 4))
sns.scatterplot(x=y_pred, y=residuals)
plt.axhline(0, color='red', linestyle='--')
plt.xlabel("Predicted Values")
plt.ylabel("Residuals")
plt.title("Residuals vs Predicted (Linearity & Homoscedasticity)")
plt.show()

# 2. Independence (Durbin-Watson)
dw = durbin_watson(residuals)
print(f"□ Durbin-Watson Statistic: {dw:.3f} (≈2 is ideal)")

# 3. Homoscedasticity (Breusch-Pagan Test)
X_test_const = sm.add_constant(X_test_scaled)
bp_test = het_breuschpagan(residuals, X_test_const)
labels = ['LM stat', 'LM p-value', 'F stat', 'F p-value']
print("\n□ Breusch-Pagan Test Results:")
print(dict(zip(labels, bp_test)))


# 4. Normality of Residuals
# Histogram
sns.histplot(residuals, kde=True)
```

```
plt.title("Histogram of Residuals (Normality Check)")
plt.show()

# Q-Q Plot
sm.qqplot(residuals, line='s')
plt.title("Q-Q Plot of Residuals")
plt.show()

# Shapiro-Wilk Test
shapiro_stat, shapiro_p = shapiro(residuals)
print(f"\n❑ Shapiro-Wilk Test p-value: {shapiro_p:.4f} (p > 0.05 ⇒ Normal)")

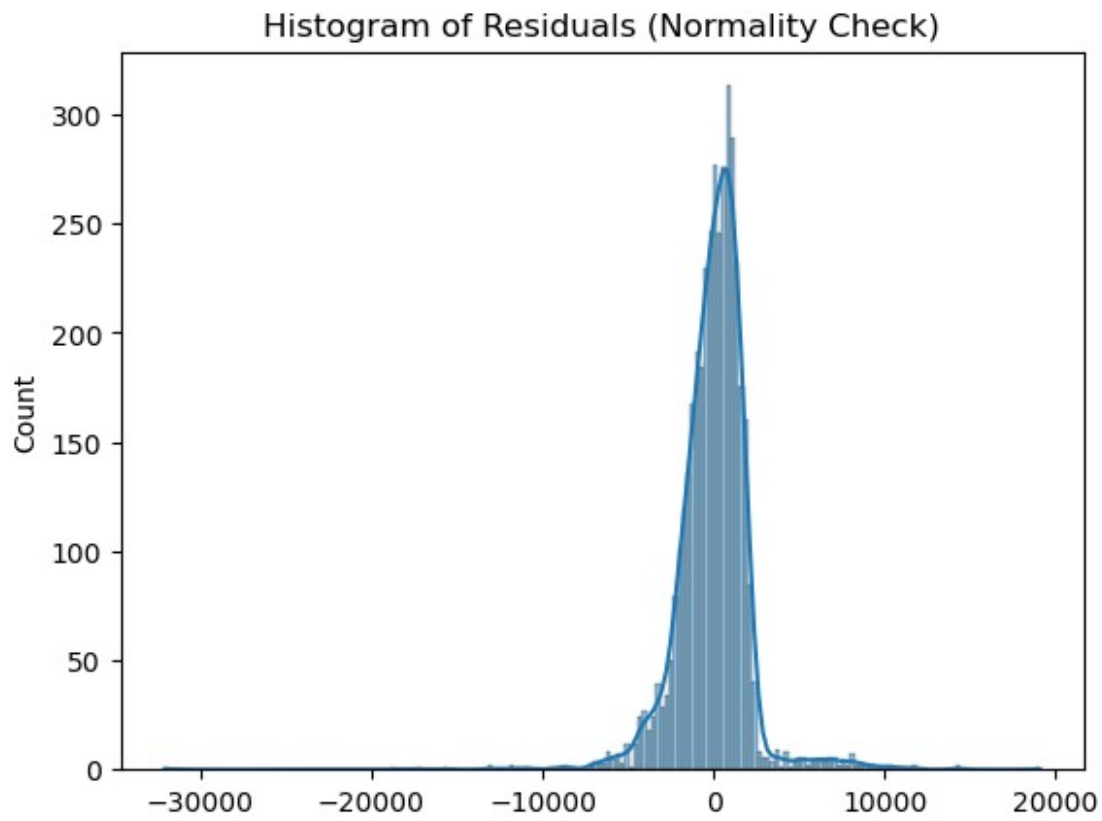
# 5. Multicollinearity (VIF)
vif_data = pd.DataFrame()
vif_data["Feature"] = [f"X{i+1}" for i in range(X_train_scaled.shape[1])]
vif_data["VIF"] = [variance_inflation_factor(X_train_scaled, i) for i in range(X_train_scaled.shape[1])]
print("\n❑ Variance Inflation Factor (VIF):")
print(vif_data)
```




```
❑ Durbin-Watson Statistic: 2.003 (≈2 is ideal)
❑ Breusch-Pagan Test Results:
```



```
{'LM stat': 1041.7415252144006, 'LM p-value': 1.4562785548263177e-219,  
'F stat': 175.67728033940617, 'F p-value': 8.205837421306955e-255}
```





□ Shapiro-Wilk Test p-value: 0.0000 ($p > 0.05 \Rightarrow$ Normal)

□ Variance Inflation Factor (VIF):

Feature	VIF
0 X1	1.208422
1 X2	1.498286
2 X3	1.496220
3 X4	1.496323
4 X5	36.984607
5 X6	1.000397
6 X7	1.031433
7 X8	37.198792

```
import pandas as pd
from statsmodels.stats.outliers_influence import
variance_inflation_factor
```

```
# Select only the first 6 features
X_train_top6 = X_train_scaled[:, 0:6]
```

```
# VIF calculation
vif_data = pd.DataFrame()
```

```
vif_data["Feature"] = [f"X{i+1}" for i in range(6)]  
vif_data["VIF"] = [variance_inflation_factor(X_train_top6, i) for i in  
range(6)]
```

```
print("\n VIF for First 6 Variables:")  
print(vif_data)
```

```
\n VIF for First 6 Variables:
```

	Feature	VIF
0	X1	1.000197
1	X2	1.498282
2	X3	1.496084
3	X4	1.496309
4	X5	1.000348
5	X6	1.000313